

Docling Quickstart: End-to-End Conversion Example

This document combines several content types on a single page — narrative text, a code sample, a data table, a bulleted list, and an embedded chart — to exercise layout segmentation and reading-order detection.

Installation

Docling is an open-source toolkit for parsing and converting documents into structured, machine-readable formats. It supports PDF, DOCX, PPTX, images, and HTML inputs, and produces a unified document representation that downstream retrieval-augmented generation (RAG) and information-extraction pipelines can consume directly. Docling is an open-source toolkit for parsing and converting documents into structured, machine-readable formats. It supports PDF, DOCX, PPTX, images, and HTML inputs, and produces a unified document representation that downstream retrieval-augmented generation (RAG) and information-extraction pipelines can consume directly.

Minimal Usage

```
from docling.document_converter import DocumentConverter

converter = DocumentConverter()
result = converter.convert("sample.pdf")
print(result.document.export_to_markdown())
```

Supported Formats

Format	Extension	OCR Required	Notes
Portable Document	.pdf	Sometimes	Native or scanned
Word Document	.docx	No	Full style fidelity
PowerPoint	.pptx	No	Slide notes preserved
Image	.png / .jpg	Yes	Always OCR'd

Deployment Checklist

- Create an OpenShift AI project and workbench.
- Deploy the docling-serve model server.
- Expose the inference route and test with a sample PDF.
- Wire the endpoint into your ingestion pipeline.

Benchmark Trend

Conversion latency (ms) over 10 runs

